

Sriram Seelamneni

Applied AI Engineer | Forward-Deployed Engineer
srirams.me | github.com/sriramsme | linkedin.com/in/sriramsme

Summary

Applied AI and Forward-deployed engineer with hands-on customer-facing delivery experience. Works directly with operators and domain experts to turn ambiguous workflows into full-stack production systems across web, mobile, backend, AI agents, and cloud infrastructure.

Professional Experience

Forward-Deployed Full-Stack Engineer

Airlink Flight School | 2025 - current

- Led technical discovery directly with the flight-school owner, instructors/pilots, and administrators, mapping how they used FlightCircle alongside paper records, manual scanning, file storage, and email-based FAA submissions.
- Owned end-to-end delivery of a web and mobile operations platform covering admissions, scheduling, grading, lesson planning, and reporting for approximately **40–50 recurring users** across multiple roles.
- Digitized the recurring student-grading process, eliminating paper grade sheets, manual scanning, record lookup, and email-based submissions; reduced administrative workload by approximately **60–70%**.
- Designed an offline-first grading workflow for instructors operating in-flight on unreliable connections, using local persistence and conflict-aware synchronization to preserve training records across devices.
- Built the initial platform with **Python, FastAPI, Firebase, and GCP**, then re-architected and migrated it to **TanStack Start, PostgreSQL, Better Auth, Neon, Railway, and Cloudflare** as the product's data, authentication, deployment, and security requirements matured.
- Owned the production delivery cycle across application architecture, data migration, authentication, deployment, Cloudflare WAF, and Turnstile.

Applied AI Engineer

allwhile.co | 2026 - present

- Conducted discovery with **10+ independent restaurant owners**, including owners operating without intelligence tools because of cost and time constraints, and owners dissatisfied with narrow, single-platform insights.
- Designed and built Allwhile as an agentic “invisible teammate” that automates competitor discovery and monitoring, business-performance and financial tracking, campaign follow-up, and local opportunity detection around events and demand signals.
- Built **Zen mode**, a low-overhead conversational interface that lets owners ask Allwhile for business and market intelligence without navigating traditional dashboards.
- Translate repeated operator pain points into product workflows, data models, agent behavior, reviewable actions, and full-stack product iterations.

Software Engineer

One Community Global | 2024-2025

- Contributed as a MERN (React.js/MongoDB) full-stack developer to build open-source Highest Good Network software.
- Built dashboard components and a supplier performance graph, fixed role-based permission leaks, and improved frontend performance by **20–30%** through caching, redundant API-call removal, and code optimization; reviewed **40+ pull requests**, authored **15+**, and helped coordinate a seven-person engineering team.

Info-Tech Assistant

Computational Mathematics, Science & Eng dept, MSU | 2023-2024

- Automated departmental data management — replaced manual Excel-based workflows with a Python/SQLite pipeline aggregating data across multiple sources, reducing administrative overhead significantly.
- Built publication tracking and reporting tooling by integrating ORCID and Zotero APIs, automated annual metrics reports covering faculty research output, funding, and publication volume
- Facilitated the transition from an outdated format to a new, optimized structure, enhancing data accessibility and relevance for departmental use

Projects

OnlyAgents: AI agent runtime

Go, WebSockets, SQLite, LLMs, Event Bus, Tool Calling | github.com/OnlyAgents

- Built an open-source, model-agnostic runtime for autonomous AI agents with executive-worker orchestration, event-driven

coordination, tool calling, connectors, and persistent memory.

- Designed a kernel-mediated execution model with a typed event bus, preventing direct agent-to-agent calls and centralizing lifecycle management, routing, and capability boundaries.
- Implemented dynamic tool gating and on-demand tool groups so agents load only the capabilities relevant to the current task, reducing unnecessary context overhead and tightening permission boundaries.
- Built layered memory across working context, semantic episode retrieval, time-scoped knowledge-graph relationships, and behavioral pattern extraction.
- Optimized the runtime into a **36 MB** single binary with approximately **27 MB idle memory** and **31 ms startup time**, enabling lightweight self-hosted deployment.

Pickle: remote dev lab for AI agents

Go, tmux, WebSockets, Tailscale, Docker, PWA | [github/pickle](#)

- Built a Go-based remote development environment exposing persistent tmux sessions, shells, local projects, development servers, and Docker Compose services through a browser and mobile-friendly interface.
- Added private remote access through Tailscale Serve and a mobile home-screen experience while keeping development sessions and project data on the host.
- Added process-aware support for Codex, Claude Code, OpenCode, Pi, and Hermes, with opt-in notifications for completed work and required approvals.

Text-to-Action: Natural Language Action System

Python, Transformers, PyTorch, Semantic Search, NER | [github/text_to_action](#)

- Built an open-source system that translates natural-language requests into executable application actions across APIs, chatbots, and voice-driven workflows.
- Implemented semantic retrieval and parameter-extraction pipelines using embeddings, LLMs, and named-entity recognition; reached **1,500+ installs**.

Enhanced Video Assistant — TechSmith Capstone

Azure, FFmpeg, Machine Learning, Python, FastAPI, React, Docker | [eva/techsmith](#)

- Built and deployed an award-winning full-stack video assistant that combined transcript analysis, metadata extraction, audio processing, and visual scene understanding.
- Integrated open-source machine-learning models to automate audio enhancement, generate multi-layered feedback, and identify key moments in video.

Picko: Personalized Recommendation System

Python, FastAPI, SQLAlchemy, PostgreSQL, React Native, Semantic Search | [picko.app](#)

- Built and deployed a recommendation engine using TF-IDF, cosine similarity, and semantic search for personalized movie, television, and book recommendations.
- Designed a **300,000+ item** relational dataset with high-concurrency query paths and automated ingestion and refresh pipelines.

Skills

Languages: Python, Go, C/C++, TypeScript, SQL

AI/ML: PyTorch, LLM applications, agent orchestration, semantic retrieval, embeddings, persistent memory, scikit-learn, NLP, computer vision, recommendation systems

Backend & Data: FastAPI, Node.js, NextJS/TanStack Start, REST APIs, OAuth, PostgreSQL, Firestore, MongoDB, SQLite

Cloud & Platform: Docker, GCP, Firebase, Neon, Railway, Cloudflare, Azure, Supabase, CI/CD, Tailscale

Frontend & Mobile: React, React Native, Expo, Tailwind CSS

Education

Michigan State University, College of Engineering | 2021-2024

Bachelor of Science, Computer Science

GPA: 3.84

Coursework: Computational Mathematics, Science & Engineering (CMSE) minor, Business cognate

Awards and Achievements

- **Auto-Owners Exposition Award** - Enhanced Video Assistant, April 2024
- **Dean's Scholarship** - \$75,000 merit-based scholarship, 2021-2024
- **Best Emerging Technology Award** - Brainwaive, January 2023